

# Serverless EV ETL Pipeline

A serverless, event-driven data pipeline for processing Electric Vehicle (EV) charging session data using AWS services. This project demonstrates how to build a cost-efficient and scalable ETL (Extract, Transform, Load) workflow using serverless technologies and Infrastructure as Code (IaC) via Terraform.

## Architecture Overview

Component	Description
Amazon S3 (Landing Zone)	Stores raw CSV files containing EV session data uploaded from data sources.
AWS Lambda (Trigger)	Automatically triggered upon file upload in S3; initiates the ETL process.
AWS Glue (Transformation Layer)	Cleans, validates, and transforms the raw data into partitioned Parquet format.
Amazon S3 (Curated Zone)	Stores transformed Parquet data in a structured, query-friendly format.
Amazon Athena (Query Layer)	Enables SQL-based querying and analysis of curated EV data directly from S3.
Terraform	Manages and provisions all infrastructure resources in a reproducible, declarative manner.

## Key Features

- Serverless and event-driven architecture — no manual triggers required.
- Infrastructure as Code (IaC) using Terraform.
- Cost-efficient — only pay for compute time when jobs run.
- Scalable — automatically adapts to data volume.
- Domain-specific design — optimized for EV charging session datasets.

## Evaluation and Review

### Strengths

- Clear architecture using AWS serverless services.
- Infrastructure as Code with Terraform for reproducibility.
- Serverless, scalable, and cost-efficient design.
- Educational and domain-focused implementation.

### Limitations

- Documentation lacks deep operational details.
- No integrated monitoring or alerting (CloudWatch).
- No schema evolution or data validation pipeline.
- Missing IAM security and encryption best practices.
- No automated testing or CI/CD workflow.

## Verdict

This project is a strong proof-of-concept implementation of a serverless ETL pipeline for EV data processing. It effectively demonstrates an end-to-end workflow using AWS-native tools. To make it production-ready, enhancements in monitoring, schema evolution, security, and testing are recommended.

## Future Enhancements

- Add schema evolution and data validation.
- Integrate CloudWatch metrics and alerts.
- Include CI/CD for automated deployments.
- Optimize Glue job partitioning for performance.
- Add IAM least privilege and encryption policies.
- Implement cost monitoring and lifecycle management.